

TigerGraph evaluation (vs Janusgraph)

Test environment (single node)

CPU	Intel(R) Xeon(R) E5-2682, 64 cores
Memory	512GB
Disk	6TB SSD
Network	25Gb/s

Datasets

Dataset	# vertices (million)	# edges (million)	Raw size (GB)
Twitter_rv	41.6	1,470	26.0
LDBC(10W)	43.0	98	6.8

1: Twitter_rv : <http://an.kaist.ac.kr/traces/WWW2010.html>

2: LDBC(10W): https://github.com/ldbc/ldbc_snb_implementations

Commands to generate the datasets (the schema can be found in the Appendix):

```
git clone https://github.com/ldbc/ldbc_snb_datagen.git
```

```
mvn clean package -DskipTests
```

```
# Edit the configuration file "params.ini" to control the amount of the data
```

```
vi params.ini
```

```
hadoop jar target/ldbc_snb_datagen-0.2.7-jar-with-dependencies.jar ./params.ini
```

Important parameters used to generate our test datasets:

- Configure the number of Persons to 100,000
- Use the following 12 tables:
 - 1) comment,
 - 2) person,
 - 3) post,
 - 4) place,
 - 5) tag,
 - 6) comment_replyOf_comment,
 - 7) comment_hasCreator_person,
 - 8) person_knows_person,
 - 9) comment_replyOf_post,

- 10) post_hashtag_tag,
- 11) post_isLocatedIn_place,
- 12) post_hasCreator_person.

Data loading

Dataset	TigerGraph (s)	Janusgraph (s)
Twitter_rv	652.09	7,982
LDBC(10W)	197.22	1,731

Size of loaded data

Dataset	Raw size (GB)	TigerGraph (GB)	Janusgraph (GB)
Twitter_rv	26.0	11	47
LDBC(10W)	6.8	13	19.2

Query latency

k-neighbor query using Twitter_rv

k-neighbor	# results (AVG)	TigerGraph (ms)	Janusgraph (ms)
1	1,611	5.60	12.50
2	6,765,143	351.00	66,700.00

LDBC Social Network Benchmark (SNB interactive)

case	TigerGraph (ms)	Janusgraph (ms)
1	9.56	3.63
2	15.99	829.63
3	12.48	301.70
4	56.13	13,816.72
5	13.33	273.58
6	54.84	6,422.00
7	10.84	195,440.00

Concluding remarks

- +3 TigerGraph loads data very fast, with low storage cost.
- +2 TigerGraph's query performance is much better than Janusgraph.

- -0 While TigerGraph's G-SQL looks somewhat like SQL, it is not. For simple/common cases, the benefit comparing to using Gremlin isn't that clear, and may make applications less portable.
- -1 TigerGraph's performance could become (much) worse when the query involves reading and/or filtering on attribute values (in addition to graph topology).
- -2 For 6 out of the 7 LDBC SNB cases, TigerGraph's performance is worse than using an open-sourced dataflow system (with our home-built Gremlin compiler); following company policy, we didn't include the actual numbers in the report.

Appendix

Janusgraph queries for LDBC SNB

1	<code>g.V().hasLabel("comment").has("id",%sL).as("x").in("comment_replyof_comment").as("y").valueMap().as("reply").select("y").out("comment_hascreator_person").valueMap().as("replycreator").select("reply","replycreator").count()</code>
2	<code>g.V().hasLabel("person").has("id",%sL).repeat(__.out("person_knows_person")).times(3).emit().dedup().has("firstname",eq("John")).count()</code>
3	<code>g.V().hasLabel("person").out("person_knows_person").as("friend").in("post_hascreator_person").has("creationdate",lt("2011-03-27")).as("post").order().by("creationdate",decr).by(id(),incr).count()</code>
4	<code>g.V().hasLabel("person").has("id",%sL).as("0").out("person_knows_person").as("x").out("person_knows_person").as("x").select(all,"x").unfold().dedup().where(neq("0")).as("friend").in("post_hascreator_person").has("creationdate",inside("2011-03-27","2011-04-27")).where(__.out("post_islocatedin_place").has("name",eq("Vietnam"))).as("post").count()</code>
5	<code>g.V().hasLabel("person").has("id",%sL).as("root").out("person_knows_person").inE("post_hascreator_person").outV().has("creationdate",inside("2011-03-27","2011-04-27")).as("post").outE("post_hastag_tag").inV().as("tag").select("tag").by("id").groupCount().order(local).by(values,decr).unfold().count()</code>

6	<code>g.V().hasLabel("person").has("id",%sL).repeat(out("person_knows_person").simplePath()).emit().times(2).dedup().in("post_hascreator_person").as("post").out("post_hashtag_tag").groupCount().by("name").order(local).by(values,decr)</code>
7	<code>g.V().hasLabel("person").has("id",%sL).in("post_hascreator_person").as("post").in("comment_replyof_post").as("comment").out("comment_hascreator_person").as("commenter").order().by(select("comment").by("creationdate")).select("commenter","comment").count()</code>

TigerGraph queries for LDBC SNB

Case 1:

```
CREATE QUERY query1(VERTEX<comments> p) FOR GRAPH ldb_10w{

    SumAccum<INT> @@cnt = 0;

    start = {p};

    reply = SELECT tgt FROM start:s-(comment_replied_comment:e)->comments:tgt;

    replyCreator      =      SELECT      tgt      FROM      reply:s-
(comment_hascreator_person:e)->person:tgt ACCUM @@cnt += 1;

    PRINT @@cnt;

}
```

Case2:

```
CREATE QUERY query2(VERTEX<person> p) FOR GRAPH ldb_10w{

    SumAccum<INT> @@cnt = 0;

    OrAccum @visited = false;

    start = {p};

    friend = {};

    friend01 = SELECT tgt FROM start:s-(person_knows_person:e)->person:tgt ACCUM
tgt.@visited += true, s.@visited += true;

    friend02 = SELECT tgt FROM friend01:s-(person_knows_person:e)->person:tgt WHERE
tgt.@visited == false ACCUM tgt.@visited += true;

    friend03 = SELECT tgt FROM friend02:s-(person_knows_person:e)->person:tgt WHERE
tgt.@visited == false ACCUM s.@visited += true;

    result = SELECT tgt FROM friend01:tgt WHERE tgt.firstName == "John" ACCUM @@cnt +=
1;
```

```

    result = SELECT tgt FROM friend02:tgt WHERE tgt.firstName == "John" ACCUM @@cnt +=
1;

    result = SELECT tgt FROM friend03:tgt WHERE tgt.firstName == "John" ACCUM @@cnt +=
1;

    PRINT @@cnt;
}

```

Case3:

```

CREATE QUERY query3(VERTEX<person> p) FOR GRAPH ldb_10w{

    SumAccum<INT> @@cnt = 0;

    start = {p};

    friend = SELECT tgt FROM start:s-(person_knows_person:e)->person:tgt;

    Post = SELECT tgt FROM friend:s-(person_create_post:e)->post:tgt WHERE
datetime_diff(tgt.creationDate, to_datetime("2011-03-27")) < 0 ORDER BY
tgt.creationDate DESC, tgt.id ASC;

    result = select tgt from Post:tgt ACCUM @@cnt += 1;

    PRINT @@cnt;
}

```

Case4:

```

CREATE QUERY query4(VERTEX<person> p) FOR GRAPH ldb_10w{

    SumAccum<INT> @@cnt = 0;

    OrAccum @visited = False;

    start = {p};

    friends={};

    friend = SELECT tgt FROM start:s-(person_knows_person:e)->person:tgt ACCUM
tgt.@visited += true, s.@visited += true;

    friendFriend = SELECT tgt FROM friend:s-(person_knows_person:e)->person:tgt WHERE
tgt.@visited == false ACCUM tgt.@visited += true;

    friends = friend UNION friendFriend;
}

```

```

    result = SELECT tgt FROM friends:s-(person_create_post:e)->post:tgt WHERE
datetime_diff(tgt.creationDate, to_datetime ("2011-03-27")) > 0 AND
datetime_diff(tgt.creationDate, to_datetime ("2011-04-27")) < 0;

    finalResult = SELECT s FROM result:s-(post_islocatedin_place:e)->place:pl WHERE
pl.name == "Vietnam" ACCUM @@cnt += 1;

    PRINT @@cnt;
}

```

Case5:

```

CREATE QUERY query5(VERTEX<person> p) FOR GRAPH ldb_10w{

    SumAccum<INT> @@cnt = 0;

    MapAccum< INT,INT> @@tagPostCounts;

    start = {p};

    friend = SELECT tgt FROM start:s-(person_knows_person:e)->person:tgt;

    Post = SELECT tgt FROM friend:s-(person_create_post:e)->post:tgt WHERE
datetime_diff(tgt.creationDate, to_datetime ("2011-03-27")) > 0 AND
datetime_diff(tgt.creationDate, to_datetime ("2011-04-27")) < 0;

    Tags = SELECT tgt FROM Post:s-(post_hastag_tag:e)->tag:tgt ACCUM @@tagPostCounts +=
(tgt.id -> 1);

    results = SELECT tgt FROM Tags:tgt ORDER BY @@tagPostCounts.get(tgt.id) DESC;

    finalresult = select tgt from results:tgt ACCUM @@cnt += 1;

    PRINT @@cnt;
}

```

Case6:

```

CREATE QUERY query6(VERTEX<person> p) FOR GRAPH ldb_10w{

    SumAccum<INT> @@cnt = 0;

    OrAccum @visited = false;

    MapAccum< STRING,INT> @@tagPostCounts;

    start = {p};

```

```

    friendSet = {};

    friend = SELECT tgt FROM start:s-(person_knows_person:e)->person:tgt ACCUM
tgt.@visited += true, s.@visited += true;

    friendFriend = SELECT tgt FROM friend:s-(person_knows_person:e)->person:tgt WHERE
tgt.@visited == false;

    friendSet = friend UNION friendFriend;

    Post = SELECT tgt FROM friendSet:s-(person_create_post:e)->post:tgt;

    Tags = SELECT tgt FROM Post:s-(post_hastag_tag:e)->tag:tgt ACCUM @@tagPostCounts +=
(tgt.name -> 1);

    results = SELECT tgt FROM Tags:tgt ORDER BY @@tagPostCounts.get(tgt.name) DESC LIMIT
10;

    PRINT results.name, @@tagPostCounts.get(results.name);
}

```

Case7:

```

CREATE QUERY query7(VERTEX<person> p) FOR GRAPH ldbc_10w{

    SumAccum<INT> @@cnt = 0;

    start = {p};

    posts = SELECT tgt FROM start:s-(person_create_post:e)->post:tgt;

    comment = SELECT tgt FROM posts:s-(post_replied_comment:e)->comments:tgt ORDER BY
tgt.creationDate ASC;

    commenter = SELECT tgt FROM comment:s-(comment_hascreator_person:e)->person:tgt
ACCUM @@cnt += 1;

    PRINT @@cnt;
}

```

Schema of the LDBC-10W

```
CREATE VERTEX person (PRIMARY_ID id UINT, id UINT, firstName STRING, lastName STRING, gender STRING, birthday STRING, creationDate STRING, locationIP STRING, browserUsed STRING)
```

```
CREATE VERTEX post (PRIMARY_ID id STRING, id STRING, imageFile STRING, creationDate STRING, locationIP STRING, browserUsed STRING, lang STRING, content STRING, length UINT)
```

```
CREATE VERTEX tag (PRIMARY_ID id UINT, id UINT, name STRING, url STRING)
```

```
CREATE VERTEX place (PRIMARY_ID id UINT, id UINT, name STRING, url STRING, tp STRING)
```

```
CREATE VERTEX comments (PRIMARY_ID id UINT, id UINT, creationDate STRING, locationIP STRING, browserUsed STRING, content STRING, length UINT)
```

```
CREATE DIRECTED EDGE person_knows_person (FROM person, TO person, creationDate STRING) WITH REVERSE_EDGE="person_known_person"
```

```
CREATE DIRECTED EDGE comment_replyof_post (FROM comments, TO post) WITH REVERSE_EDGE="post_replied_comment"
```

```
CREATE DIRECTED EDGE comment_replyof_comment (FROM comments, TO comments) WITH REVERSE_EDGE="comment_replied_comment"
```

```
CREATE DIRECTED EDGE post_hascreator_person (FROM post, TO person) WITH REVERSE_EDGE="person_create_post"
```

```
CREATE DIRECTED EDGE post_hastag_tag (FROM post, TO tag) WITH REVERSE_EDGE="tag_hadtag_post"
```

```
CREATE DIRECTED EDGE comment_hascreator_person (FROM comments, TO person) WITH REVERSE_EDGE="person_hadcreator_comment"
```

```
CREATE DIRECTED EDGE post_islocatedin_place (FROM post, TO place) WITH REVERSE_EDGE="place_waslocatedin_post"
```

```
CREATE GRAPH ldbc_10w (person, post, tag, place, comments, person_knows_person, comment_replyof_post, comment_replyof_comment, post_hascreator_person, post_hastag_tag, comment_hascreator_person, post_islocatedin_place)
```